


AI-assisted delivery, wired into Jira and Confluence.

sdhc-work guides AI coding agents through the process we already run and keeps Jira, Confluence and Bitbucket or GitHub current automatically. Less status admin, more verified, reviewed code.

THE LOOP



Available today **WORKING · TESTED**

- ✓ **Two tracks.** PRD-first for large features; straight from Jira for small work and bugs.
- ✓ **Automatic ticket filing.** A PRD becomes dependency-linked Jira stories plus a Confluence breakdown, safe to re-run.
- ✓ **Status that updates itself.** Stories move to In Progress, then Done, as work is planned and verified.
- ✓ **Evidence, not just status.** Verification reports and bug root-cause analyses land in Confluence, linked from Jira.
- ✓ **PRs and AI code review.** Bitbucket Data Center and GitHub: inline findings, merge gate, verdict in Jira.
- ✓ **Our stack, both ways.** Jira and Confluence Data Center and Cloud, each verified live.
- ✓ **Packaged and portable.** 15 reusable AI skills in one plugin, portable to other AI agents.

On the roadmap **PLANNED · IN JIRA**

NEXT · TEAM ROLLOUT

Safety and correctness fixes from an independent review

Multi-developer support, setup health check, one-command "start ticket"

Stories close on merge from CI; production-ready CI review pipelines

THEN · LEARN AND PROVE

Retrospective loop that improves our rules and AI skills

Measured skill quality against a no-AI baseline

Pilot on a live project

LATER · DIFFERENTIATE

Hands-off "work this ticket" on Data Center

Tests generated from acceptance criteria

Delivery metrics: cycle time, rework, cost per ticket

Evidence-gated releases; incident-to-bug intake

Less admin

Jira and Confluence reflect the real state of work without anyone updating them by hand.

Quality by default

Every change is verified and reviewed; every bug leaves a regression test and a prevention rule.

Built for where we are

Data Center today, Cloud-ready ahead of Atlassian's Data Center end of life in 2029.